Neural Networks

November 11, 2015

Abstract

In this short paper we will discuss the fundamentals of neural networks and their implementation in detail. We will give a general overview of how neural networks work, discuss calculation of the gradient and implementation of back-propagation, and test our results on some sample data (classifying toys into three subclasses) and a real dataset (the MNIST data for classifying pixels of handwritten letters). We will discuss our findings and analyze them in the context of choosing parameters and learning for our neural network.

1 Implementation Details

Neural networks have been around for at least a few decades, but only recently have them become popular as a method for learning parameters that can correctly translate an input into an output. This is because of increased computational power, a greater availability of training data, as well as the fact that more complex models, like deep neural nets, are actually easy to train - the same back-propagation that works to update normal neural networks works just as well for multiple hidden layers. Figure 1 shows a basic


Figure 1: 1-hidden layer neural networ taken from Bishop

neural network. The x on the left is the input, with each of the d dimensions acting as a separate node. These are then multiplied by the appropriate weights, added to a *bias* (representing x_0 here), in order to obtain the *activations* a_j for $1 \leq j \leq M$. These activations are then put through some non-linear map, in this case the sigmoid function, to obtain the values at the first hidden layer z_j. Mathematically, this looks like:

$$a_j^{(1)} = \sum_{i=1}^{d} w_{ji}^{(1)} x_i + w_{j0}^{(1)}$$

$$z_j = g(a_j^{(1)}) = \frac{1}{1 + e^{-a_j^{(1)}}}$$

To arrive at a simple one-hidden-layer neural network, we do this all over again, using our previous hidden layer node values (our z_j) instead of our inputs x_i as the inputs to the second layer. The equations for this layer are analogous:

$$a_k^{(2)} = \sum_{j=1}^{d} w_{kj}^{(2)} x_j + w_{k0}^{(1)}$$

$$f_k = g(a_k^{(2)}) = \frac{1}{1 + e^{-a_k^{(2)}}}$$

In this paper, we will consider a loss function in the form of a negative log likelihood, taking the form:

$$\ell(w) = \frac{1}{N} \sum_{i=1}^{N} \sum_{k=1}^{K} \left[-y_k^{(i)} \log(h_k(x^{(i)}, w)) - (1 - y_k^{(i)}) \log(1 - h_k(x^{(i)}, w)) \right]$$

where $f_k = h_k$ and y_k is a K-dimensional one-hot vector with all 0 values except for a 1 in the kth dimension. If we optimize this directly, however, we will often overfit to the training data (of which we have n samples from $x^{(1)}$ to $x^{(n)}$ - we distinguish these from x_i, which are the features of one particular sample that we will consider at a time). Thus, we add a regularization term on the weights $w^{(1)}$ and $w^{(2)}$, so that we try to minize:

$$J(w) = \ell(w) + \lambda(||w^{(1)}||_F^2 + ||w^{(2)}||_F^2)$$

To do this, we can use our gradient descent methods from previous examinations of regression and classification; in this scenario, we want $\nabla_{w_1} J(w)$ and $\nabla_{w_2} J(w)$, which we can calculate analytically. First we compute:

$$\frac{\partial J(w)}{\partial h_k(x^{(i)}, w)} = -\frac{y_k^{(i)}}{h_k(x^{(i)}, w)} + \frac{1 - y_k^{(i)}}{1 - h_k(x^{(i)}, w)} + 2\lambda(||w^{(1)}|| + ||w^{(2)}||)$$

From the lecture notes, we have an expression for $\nabla_{w_k^{(2)}} J(w)$:

$$\nabla_{w_k^{(2)}} J(w) = \frac{\partial J(w)}{\partial h_k(x^{(i)}, w)} (\tilde{g}'(a_k^{(2)}))\mathbf{z}$$

$$= \frac{\partial J(w)}{\partial h_k(x^{(i)}, w)} \frac{a_k^{(2)}}{e^{a_k^{(2)}} + 1} \mathbf{z}$$

and this can be implemented with gradient descent and back-propagation to minimize the error. Similarly, we can calculate $\nabla_{w_1} J(w)$. First we let

$$\frac{\partial J(w)}{\partial a_k^{(2)}} = \frac{\partial J(w)}{\partial h_k(x^{(i)}, w)} \frac{a_k^{(2)}}{e^{a_k^{(2)}} + 1} = \delta_k^{(2)}$$

so that

$$\nabla_{w_k^{(2)}} J(w) = \delta_k^{(2)} \mathbf{z}$$

Then, if we keep applying the chain rule like we did previously for $\nabla_{w_k^{(2)}} J(w)$, we'll get:

$$\nabla_{w_j^{(1)}} J(w) = \left(\frac{\partial J(w)}{\partial a_j^{(1)}} \right) \left(\nabla_{w_k^{(1)}} a_j^{(1)} \right)$$

2

where

$$\frac{\partial J(w)}{\partial a_j^{(1)}} = \sum_{k=1}^{K} \left(\frac{\partial J(w)}{\partial a_k^{(2)}} \right) \left(\frac{\partial a_k^{(2)}}{\partial a_j^{(1)}} \right)$$

$$= \sum_{k=1}^{K} \delta_k^{(2)} \cdot w_{kj}^{(2)} \cdot g'(a_j^{(1)})$$

$$= \sum_{k=1}^{K} \delta_k^{(2)} \cdot w_{kj}^{(2)} \cdot \frac{a_j^{(1)}}{e^{a_j^{(1)}} + 1} = \delta_j^{(1)}$$

When we combine everything, we now have

$$\nabla_{w_j^{(1)}} J(w) = \delta_j^{(1)} \mathbf{x}.$$

We can now use these formulas to implement our gradient descent or stochastic gradient descent based on our data and our choosing. We tested our code and implementation with several sample values.

2 Applications

First, we can investigate the use of our neural network in classifying toy data. The toys are separated into three different classes, and we are given training, validation, and test data for our model. We use cross-validation to help us optimize the number of hidden nodes, and we also test both batch and stochastic gradient descent to help us find the optimal parameters. We can initialize our parameters with $w_{ij}^{(1)}$ and $w_{jk}^{(2)}$ set randomly between 0 and 1 for all i, j, k and $M = 1, 2, 5, 10$. d and K are determined by the input vector and the number of "groups" we wish to classify into.

We aggregate our results by cross-validation - this gives us more confidence that our answer is correct by utilizing as much of the data as we can to optimize our parameters and weight matrices for classification of the toy data. Our solutions vary between batch and stochastic gradient descent. The best value of M, the number of hidden layers, we obtained was $M = 4$ with an accuracy of 99% and the best value of M we obtained with stochastic gradient descent as opposed to normal batch gradient descent was $M = 3$. The best value of λ was 0 in both cases.

In addition, we tested our model on real-life MNIST data, a well-known corpus of handwritten digits from 0-9 that are often used as a benchmark for classification algorithms. This dataset that we're working with, however, only includes points with classes 1 - 6. The whole process is the same as with the toy data, just that we're working with more parameters and many more dimensions. We varied the number of hidden nodes M from 5 to 100, the tradeoff regularization parameter λ from 0 to 1, the step size of the gradient descent between 0.0001 and 0.01, as well as experimented between normal gradient descent and stochastic gradient descent. Our choice of parameters ended up being fairly stable in producing large accuracies (around 90%, depending on the stochastic initial parameters) and matches - we tended to favor a value of M around 25, $\lambda = 0.03$, a step size of 0.005, and there was no noticeable difference between normal gradient descent and stochastic gradient descent (although SGD performed slightly better, and was significantly faster in runtime).

The choice of parameters is reasonable given the problem. The more data we have, in general the more complex the parameters will be in an attempt to fit all the data. However, this is balanced by the λ that seeks to prevent overfitting, and so restricts the size of M to be within a reasonable range.

**Figure 2:** `toy_1 lambda_0 batch`

Figure 3: `toy_1 lambda_1 batch`

**Figure 4:** `toy_1 lambda_0 sgd`

Figure 5: `toy_1 lambda_1 sgd`

Table 1: Varying M vs λ to calculate accuracy

M	λ	accuracy
5	.005	15.6%
5	.001	17.0%
10	.005	18.4%
10	.001	17.9%
20	.005	84.8%
20	.001	83.4%
25	.005	88.7%
25	.001	90.2%
40	.005	87.3%
40	.001	86.9%
100	.005	20.1%
100	.001	19.2%